

Recap: Basic Structured Query Language (SQL)

SQL is a standard language for querying and manipulating data in RDBMS

- CREATE TABLE with PRIMARY KEYS/FOREIGN KEYS
- SFW query

ICCS240 Database Management

SQL **(cont.)**

Many slides in this lecture are either from or adapted from slides provided by
Theodoros Rekatsinas, UW-Madison
Abdu Alawini. UIUC

Basic form of SQL Query

SELECT	<attributes: A_1, \dots, A_m >
FROM	<one or more relations: R_1, \dots, R_k >
WHERE	<predicate/conditions: C >

Call this a select-from-where or SFW query.

The result of an SQL query is a relation.

You may run/try examples from
sqllex_run1.sql along with the
following slides

Simple SQL Query: Selection

Selection (σ) is the operation of filtering a relation's tuples on some condition

PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

```
SELECT *  
FROM   Product  
WHERE  Category = 'Gadgets'
```

$\sigma_{Category='Gadgets'}(Product)$



PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks

Simple SQL Query: Projection

Projection (Π) is the operation of producing an output table with tuples that have a subset of their prior attributes

PName	Price	Category	Manufacturer
Gizmo	\$19.99	Gadgets	GizmoWorks
Powergizmo	\$29.99	Gadgets	GizmoWorks
SingleTouch	\$149.99	Photography	Canon
MultiTouch	\$203.99	Household	Hitachi

```
SELECT Pname, Price, Manufacturer
FROM   Product
WHERE  Category = 'Gadgets'
```

$\Pi_{PName, Price, Manufacturer}(\sigma_{Category = 'Gadgets'}(Product))$




PName	Price	Manufacturer
Gizmo	\$19.99	GizmoWorks
Powergizmo	\$29.99	GizmoWorks

SQL Query (may) create a result with a new schema

Input schema

Product(PName, Price, Category, Manufacturer)

```
SELECT Pname, Price, Manufacturer  
FROM   Product  
WHERE  Category = 'Gadgets'
```



Output schema

Answer(PName, Price, Manufacturer)

A few tips about SQL

- SQL **commands** are case *insensitive*:
 - Same: SELECT, Select, select
 - Same: Product, product
- **Values** are **not**:
 - Different: 'Seattle', 'seattle'
- Use single quotes for constants:
 - ✓ • 'abc'
 - ✗ • "abc"

LIKE: simple string pattern matching

```
SELECT *  
FROM Products  
WHERE PName LIKE '%gizmo%'
```

s **LIKE** p: pattern matching on strings


p may contain two special symbols:

% = any sequence of characters


_ = any single character

DISTINCT: eliminate duplicates

```
SELECT DISTINCT Category  
FROM Product
```



Category
Gadgets
Photography
Household



```
SELECT Category  
FROM Product
```



Category
Gadgets
Gadgets
Photography
Household

ORDER BY: sorting the results

```
SELECT    PName, Price, Manufacturer
FROM      Product
WHERE     Category='gizmo' AND Price > 50
ORDER BY Price, PName
```

Ties are broken by the second attribute on the ORDER BY list, etc.

Ordering is ascending, unless you specify the DESC keyword.

NULL VALUES

VALUES CAN BE *NULL*

- Tuples in SQL relations can have *NULL* as a value for one or more components.

- *Comparison:*

When any value is compared with *NULL*, the truth value is *UNKNOWN*.

- *Missing value*: e.g., we know 'Canon' company has some address, but we don't know what it is.
- *Inapplicable*: e.g., the value of attribute *spouse* for an unmarried person.

Comparison of NULLs and values

- The logic of conditions in SQL is really 3-valued logic:

TRUE, FALSE, UNKNOWN

- *Comparison:*

When any value is compared with NULL, the truth value is **UNKNOWN**

- *Outcome:*

A query only produces a tuple in the answer

if its truth value for the WHERE clause is TRUE (not FALSE or UNKNOWN)

Three-Valued Logic

- Think of TRUE = 1, FALSE = 0, and UNKNOWN = $\frac{1}{2}$
- AND = MIN(); OR = MAX(), NOT(x) = $1-x$
- Example:

TRUE **AND** (FALSE **OR** NOT(UNKNOWN))
=

Try **SELECT TRUE AND (FALSE OR NOT(NULL)) ;**

Tricky Example!

Sell	bar	beer	price
	Joe	Asahi	NULL

SELECT
FROM
WHERE

bar
Sell
price < 2.00 OR price >= 2.00

← UNKNOWN → ← UNKNOWN →

← UNKNOWN →

Surprising Result?

2-Valued Laws \neq 3-Valued Laws

Some common laws, like the commutativity of AND, hold in 3-valued logic.

$$p \text{ AND } q \quad \equiv \quad q \text{ AND } p$$

But many others do not!

Example: (complement?)

In 2-valued law: $p \text{ OR NOT } p \quad \equiv \quad \text{TRUE}$

In 3-valued law: When $p = \text{UNKNOWN}$,
L.H.S. is $\text{MAX}(\frac{1}{2}, (1 - \frac{1}{2})) = \frac{1}{2} = \text{UNKNOWN}$

Leverage relevant issue by testing for NULL

Instead of relying on there not being NULLs or getting the default behavior for NULLs, can override this.

Test for NULL explicitly:

x IS NULL

x IS NOT NULL

```
SELECT bar
FROM Sell
WHERE price < 2.00 OR price >= 2.00
      OR price IS NULL
```


Multi-Relation query

Foreign Key Constraints

Suppose we have the following schema:

```
Students(sid: string, name: string, gpa: float)
Enrolled(student_id: string, cid: string, grade: string)
```

Students			Enrolled		
sid	name	gpa	student_id	cid	grade
101	Bob	3.2	123	564	A
123	Mary	3.8	123	537	A+



Q: *student_id alone is not a key- what is?*

And we want to impose the following constraint:

'Only bona fide students may enroll in courses' i.e. a student must appear in the Students table to enroll in a class

We say that student_id is a foreign key that refers to Students

Declaring foreign keys

```
Students(sid: string, name: string, gpa: float)  
Enrolled(student_id: string, cid: string, grade: string)
```

```
CREATE TABLE Enrolled(  
    student_id CHAR(20),  
    cid CHAR(20),  
    grade CHAR(10),  
    PRIMARY KEY (student_id, cid),  
    FOREIGN KEY (student_id) REFERENCES Students(sid)  
)
```

Foreign keys & Update operations

```
Students(sid: string, name: string, gpa: float)
```

```
Enrolled(student_id: string, cid: string, grade: string)
```

What if we insert a tuple into Enrolled, but no corresponding student?

INSERT is rejected (foreign keys are constraints)!

What if we delete a student?

- | | | |
|---|-------------------------|----------|
| 1. Disallow the delete | (ON DELETE NO ACTION) | DEFAULT! |
| 2. Remove all of the courses for that student | (ON DELETE CASCADE) | |
| 3. <i>SQL allows via NULL</i> | (ON DELETE SET NULL) | |
| 4. SQL allows via default values | (ON DELETE SET DEFAULT) | |

(similar idea for update?)

JOINS of multiple relations

Product	(pname, price, category, maker)
Company	(cname, stock, country)
Purchase	(buyer, seller, store, product)
Person	(pname, phoneNumber, city)

Find all products under \$100 manufactured in USA; return their name, price and its maker.

Several equivalent ways to write a basic join in SQL:

A join between tables returns all unique combinations of their tuples which meet some specified join condition

```
SELECT pname, price, maker
FROM   Product, Company
WHERE  maker = cname
      AND country = 'USA'
      AND price <= $100
```

```
SELECT pname, price, maker
FROM   Product JOIN Company
      ON maker = cname
WHERE  price <= $100 AND country = 'USA'
```

Example

Product2

PName	Price	Category	Manuf
Gizmo	\$119	Gadgets	GWorks
PowerGizmo	\$29	Gadgets	GWorks
SingleTouch	\$49	Photography	Micron
MultiTouch	\$203	Household	Hitachi

Company2

Cname	Stock	Country
GWorks	25	USA
Micron	65	USA
Hitachi	15	Japan



```
SELECT pname, price, manu  
FROM   Product2, Company2  
WHERE  manu = cname  
        AND  country = 'USA'  
        AND  price <= $100
```

PName	Price	Maker
PowerGizmo	\$29	Gworld
SingleTouch	\$49	Micron

Ambiguity in Multi-Relation

```
Person(name, address, worksfor)
Company(name, address)
```

```
SELECT DISTINCT name, address
FROM           Person, Company
WHERE          worksfor = name
```

Which “address” does this refer to?

Which “name”s??

Disambiguating in Multi-Relation

Person(name, address, worksfor)

Company(name, address)

Both equivalent
ways to resolve
variable
ambiguity

```
SELECT DISTINCT Person.name, Person.address
FROM             Person, Company
WHERE            Person.worksfor = Company.name
```

```
SELECT DISTINCT p.name, p.address
FROM             Person p, Company c
WHERE            p.worksfor = c.name
```

Tip: Always prefix with relation name to make it clear/easier to read. But optional!


Naming Convention

- To avoid ambiguities, every attribute name has two components

Relation . Attribute

- When there is no ambiguity, one can drop the relation name component

```
SELECT  person.name, person.income
FROM    person
WHERE   person.age < 30
```



```
SELECT  name, income
FROM    person
WHERE   age < 30
```

Your playtime:

Download and run **sqlex_data2.sql**

From the dataset:

- From the table **person**, compute a new table by selecting only the persons with an income between 20 and 30, and adding an attribute that has, for every tuple, the same value as **income**.
- The fathers of persons who earn more than 20K
- Father and mother of every person
- Persons that earn more than their father, showing name, income, and income of the father

MotherChild

mother	child
Lisa	Mary
Lisa	Greg
Anne	Kim
Anne	Phil
Mary	Andy
Mary	Rob

FatherChild

father	child
Steve	Frank
Greg	Kim
Greg	Phil
Frank	Andy
Frank	Rob

Data from Werner Nutt

Person		
name	age	income
Andy	27	21
Rob	25	15
Mary	55	42
Anne	50	35
Phil	26	30
Greg	50	40
Frank	60	20
Kim	30	41
Mike	85	35
Lisa	75	87

Operational Semantics

(Similar to all single-relation queries)

1. Start with the product of all the relations in the FROM clause.
2. Apply the selection condition from the WHERE clause.
3. Project onto the list of attributes and expressions in the SELECT clause.

Meaning (Semantics) of SQL Queries

SELECT	A_1, A_2, \dots, A_k
FROM	R_1, R_2, \dots, R_n
WHERE	conditions

Translate to Relational Algebra

$$\Pi_{A_1, A_2, \dots, A_k} \left(\sigma_{\text{conditions}} (R_1 \times R_2 \times \dots \times R_n) \right)$$

Select-From-Where queries are precisely Join-Select-Project

Meaning (Semantics) of SQL Queries

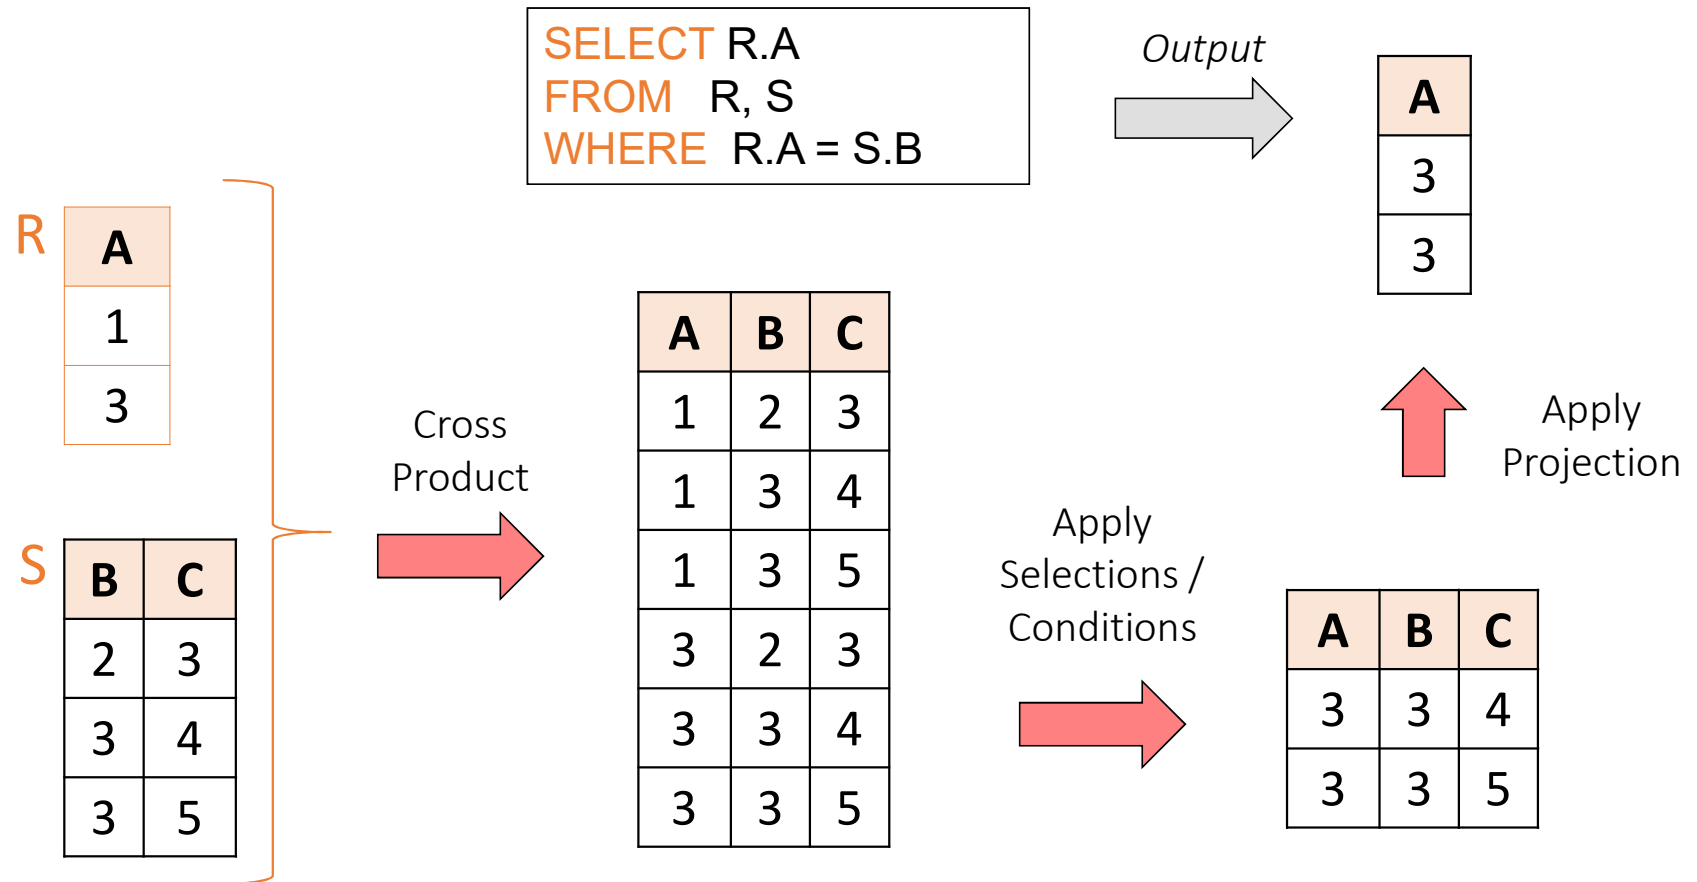
SELECT	A_1, A_2, \dots, A_k
FROM	R_1, R_2, \dots, R_n
WHERE	conditions

Nested Loop

```
Answer = {}  
for x1 in R1 do  
    for x2 in R2 do  
        ...  
        for xn in Rn do  
            if Conditions  
                then Answer = Answer  $\cup$  { (A1,...,Ak) }  
return Answer
```

Almost never the *fastest* way to compute it!

An example of SQL semantics



Operational Semantics: more explanation

```
SELECT R.A  
FROM   R, S  
WHERE  R.A = S.B
```

1. (From) Take **cross product**:

$$X = R \times S$$

Cross product ($A \times B$) is the set of all unique tuples in A, B

Ex: $\{a,b,c\} \times \{1,2\}$
 $= \{(a,1), (a,2), (b,1), (b,2), (c,1), (c,2)\}$

2. (Where) Apply **selections / conditions**: = Filtering!

$$Y = \{(r,s) \in X \mid r.A == r.B\}$$

3. (Select) Apply **projections** to get final output: = Returning only *some* attributes


$$Z = (y.A,) \text{ for } y \in Y$$

Remembering this order is critical to understanding the output of certain queries

An unintuitive query

```
SELECT DISTINCT R.A  
FROM   R, S, T  
WHERE  R.A=S.A OR R.A=T.A
```

What does it compute?



But what if $S = \emptyset$?

Go back to the semantics!