

ICCS240 Database Management

Storage and File Structure

Many slides in this lecture are either from or adapted from:

slides provided by Kazuhiro Minami, UIUC

texts from Database Management Systems by R. Ramakrishnan and J. Gehrke

Two Perspectives on DBMS

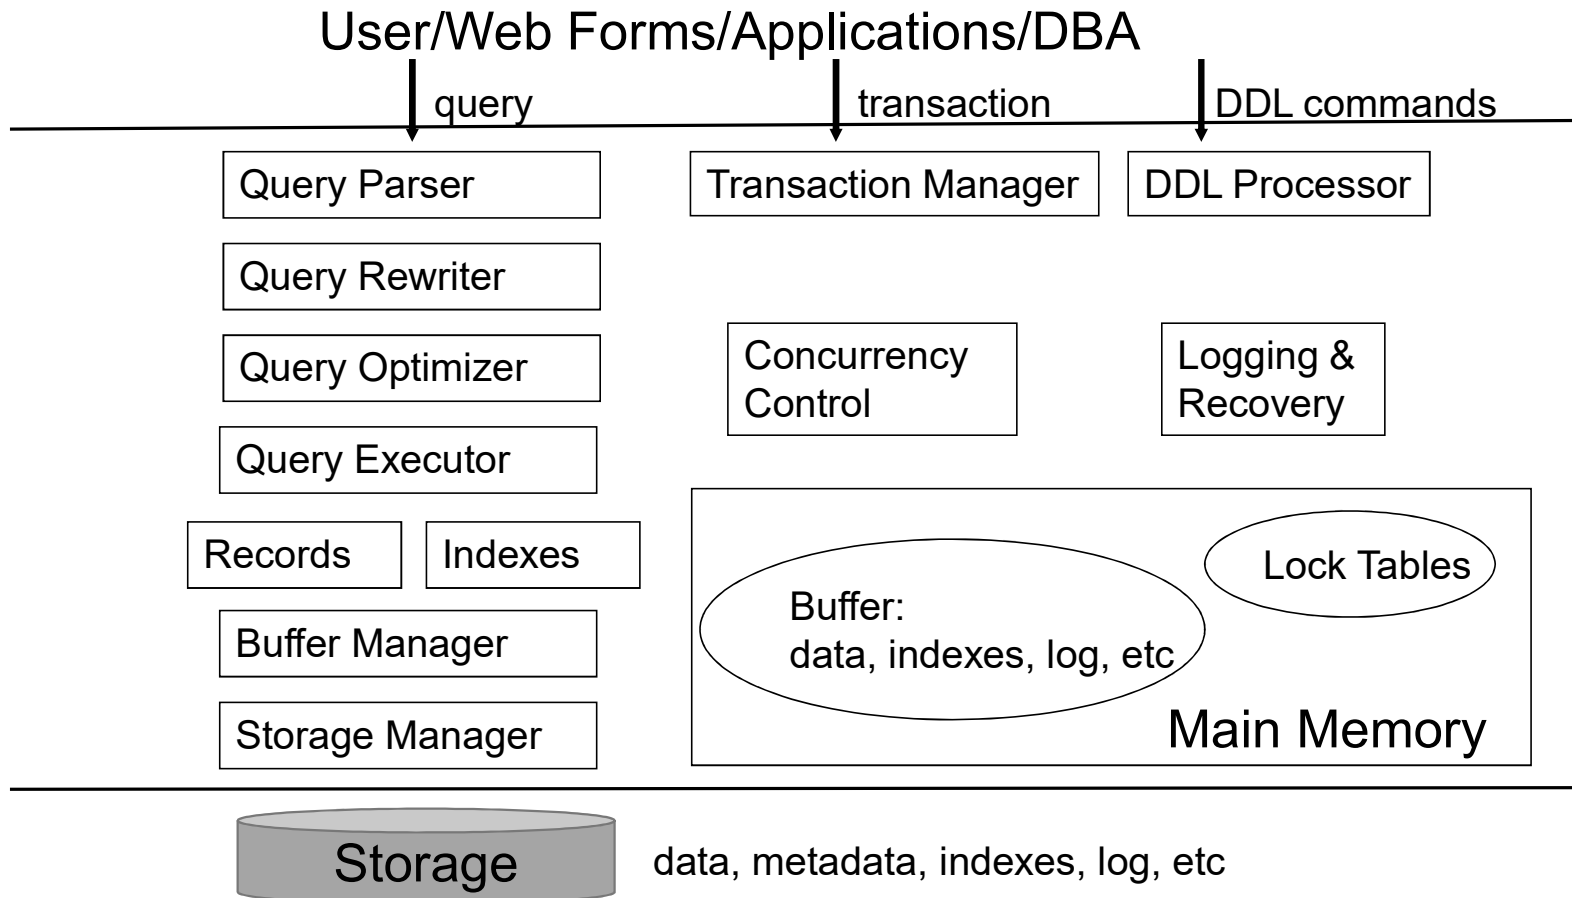
- **User perspective**

- How to use a database system
 - Database design
 - Database programming

- **System perspective**

- How to design and implement a database system
 - Storage management
 - Query processing
 - Transaction management

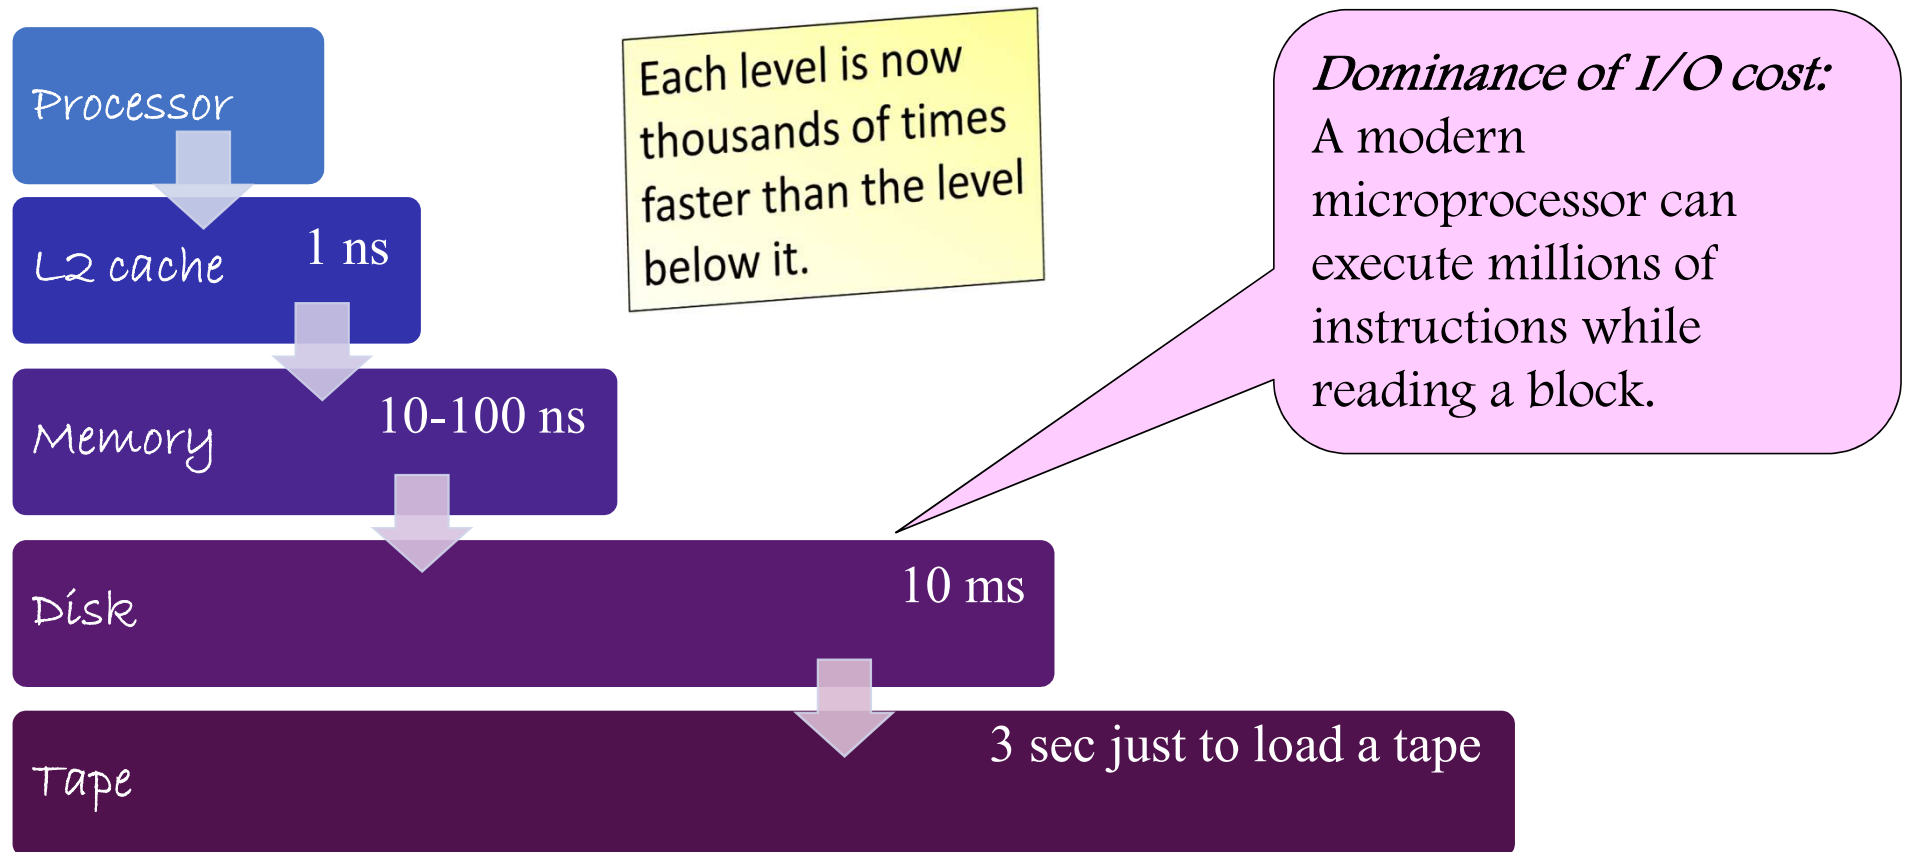
The Big Picture — DBMS Architecture



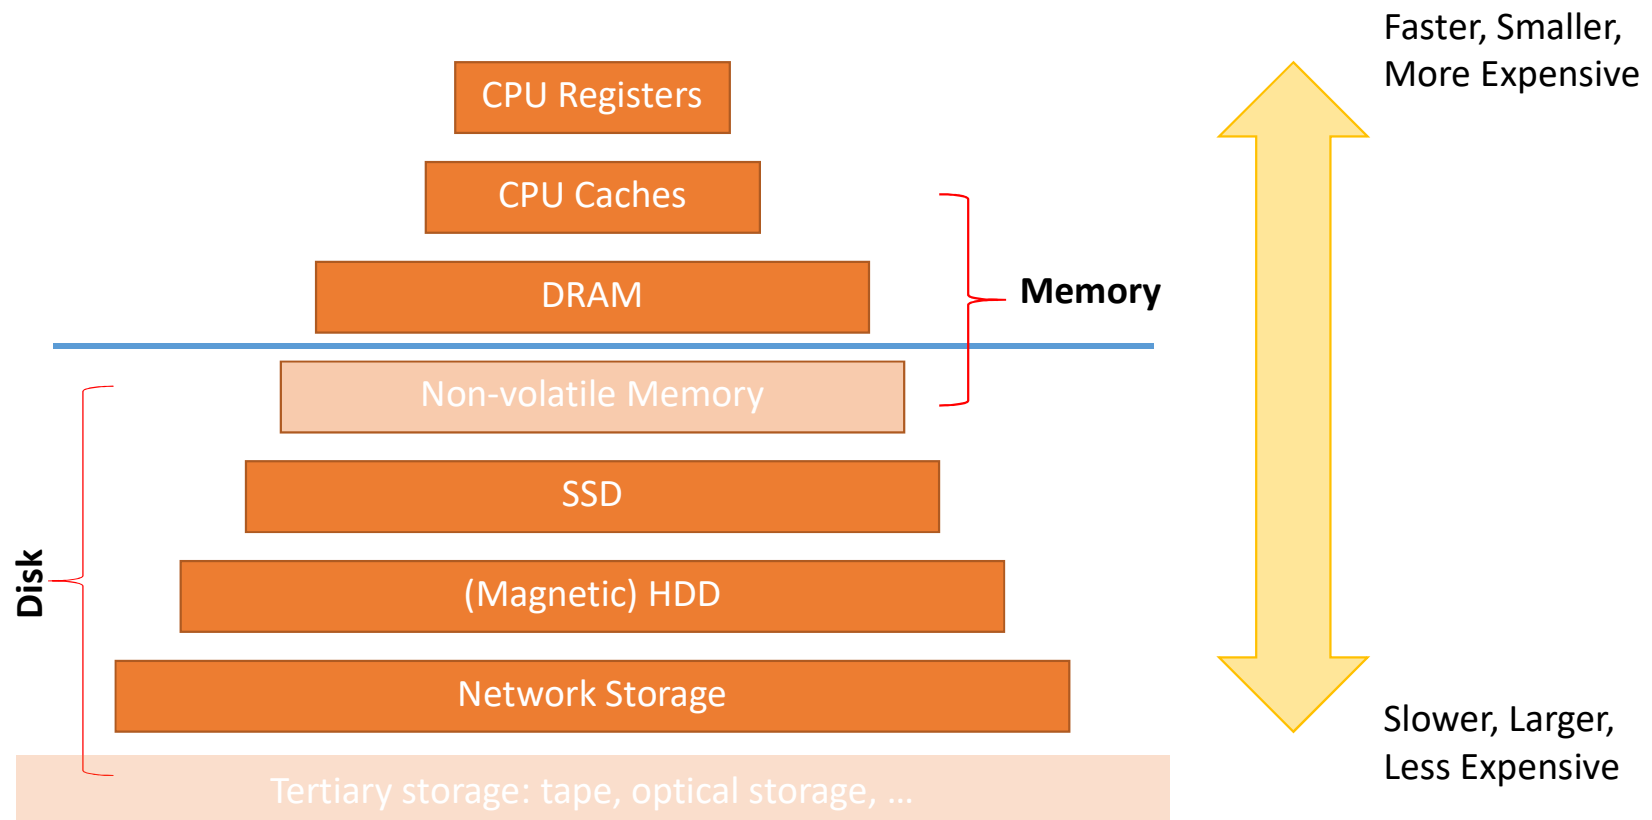
Access Times

0.5 ns	L1 Cache Ref	← 0.5 sec
7 ns	L2 Cache Ref	← 7 sec
100 ns	DRAM	← 100 sec
150,000 ns	SSD	← 1.7 days
10,000,000 ns	HDD	← 16.5 weeks
~30,000,000 ns	Network Storage	← 11.4 months
1,000,000,000 ns	Tape Archives	← 31.7 years

The Memory Hierarchy



Storage Pyramid – Where should we store data?




Goals

Allow the DBMS to manage databases that exceed the amount of (fast) memory available.

But reading/writing to disk is expensive, so must manage it carefully
—to minimize stalls and performance degradation.

Hard-drive Mechanism

Good for long sequential reads; bad for seek-heavy workloads



- Magnetic disks support **direct access** to a desired location.
- Data is stored on disk in units called **disk blocks**, which is the unit of reading or writing.
 - Size of disk block can be set.
 - Typically, 4K, 8K, 16K
- Surface of platter divided into circular **tracks**.
 - Over 50K-100K tracks per platter
- Each track is divided into **sectors**
 - A sector is the smallest unit of data that can be read or written.
 - Sector size typically 512 B
 - Typically, 500-1000 sectors/track $\approx 10^5$ bytes/track
- **Platter** may have one or two surfaces
- **Cylinder** = set of all tracks with the same radius
 - Cylinder i^{th} consists of i^{th} track of all the platters

DBMS is primarily optimized for HDD

— implications:

- Seek time and rotational delay dominate!
- Key to lower I/O cost: **reduce seek/rotation delays!**
- How to minimize seek and rotational delays?
 - Blocks on same track, followed by
 - Blocks on same cylinder, followed by
 - Blocks on adjacent cylinder
 - Hence, sequential arrangement of blocks of a file is a big win!

Disk Organization: Reliability & Performance

Many Disks vs. One Disk


Although disks provide cheap, non-volatile storage for DBMSs, they are usually bottlenecks for DBMSs.

Factors: **Reliability**, **Performance**, **Capacity**

How about adopting multiple disks?

- More data can be held as opposed to one disk **Capacity!**
- Data can be stored redundantly; hence, if one disk fails, data can be found on another **Reliability!**
- Data can be accessed concurrently **Performance!**

Logical Volume Managers (LVMs)



LVMs give appearance of a single logical disk.

- **SPANNING:**

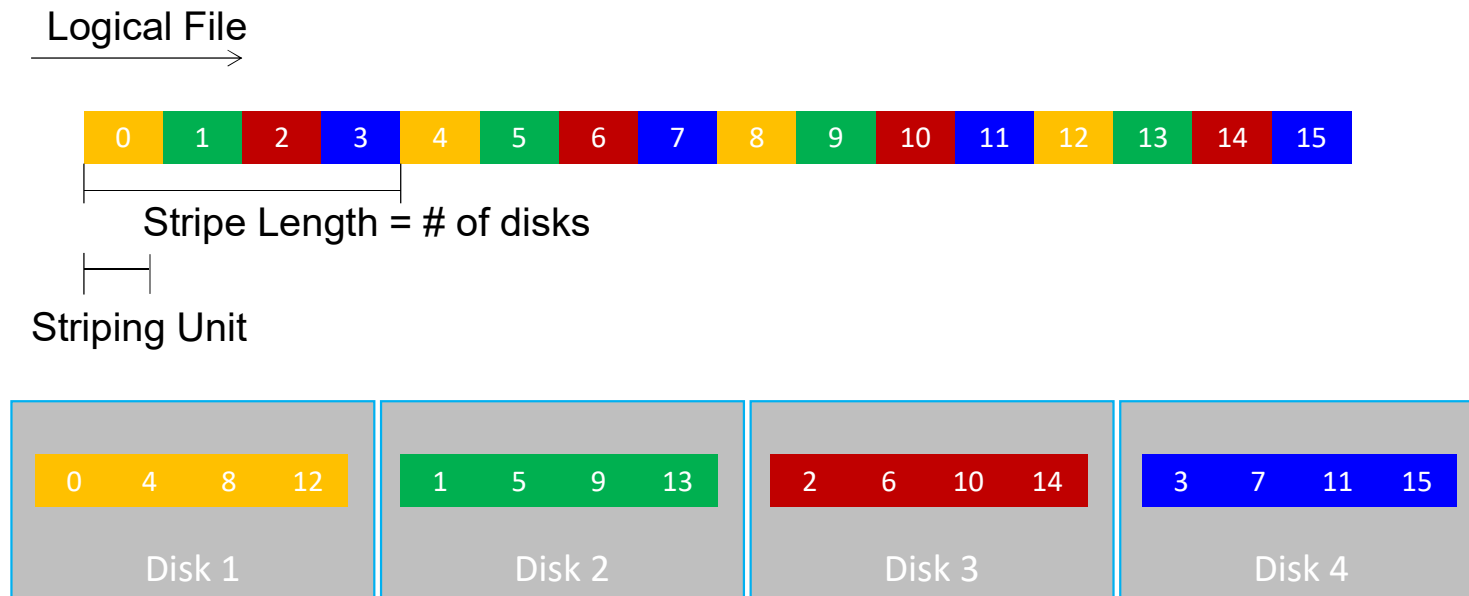
- LVM transparently maps a larger address space to different disks

- **MIRRORING:**

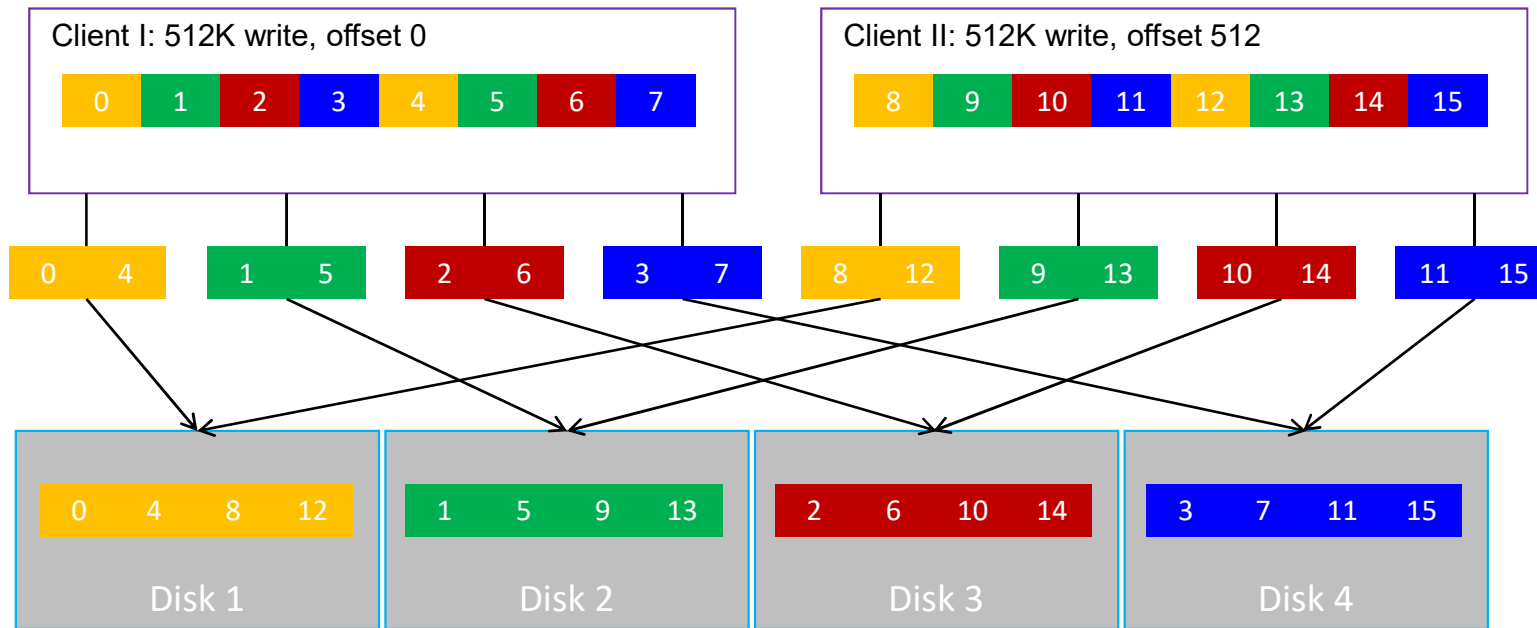
- Each disk can hold a separate, identical copy of data
- LVM directs writes to the same block address on each disk
- LVM directs a read to any disk (e.g., to the less busy one)

Data Striping

To achieve parallel accesses, we can use a technique called data striping



Data Striping (cont.)



Trade-off:

- Small striping unit values: higher parallelism, higher disk seek and rotational delays
- Large striping unit values: lower parallelism, decrease disk seek and rotational delays

RAID: Redundant arrays on independent disks


A standard techniques to improve performance and reliability when multiple disks are available.

- Improved reliability by redundant storage of data
- Reduced access cost by exploiting parallelism

(although there is obviously a trade-off between increased capacity and increased reliability via redundancy)


RAID combines mirroring and striping!

RAID0 – Striping + No Fault Tolerance



RAID1 – Striping + Mirroring

Improve READs; Impair WRITEs



RAID2-6 ...

The higher levels of raid incorporate various combinations of block/bit-level striping, mirroring, and error correcting codes

The differences are primarily in

- the kind of error checking/correcting codes that are used
- where error correcting codes parity bits are stored

Disk Management

Disk Space Management

- **Disk space manager** manages space on disk
- Disk space manager supports the concept of a **page** as a unit of data
- The size of a page is chosen to be the size of a disk block
- Useful to allocate a **sequence of pages as a contiguous sequence of blocks** to hold data frequently accessed.
- Disk space manager *keeps track* of blocks in use or location of blocks on disk

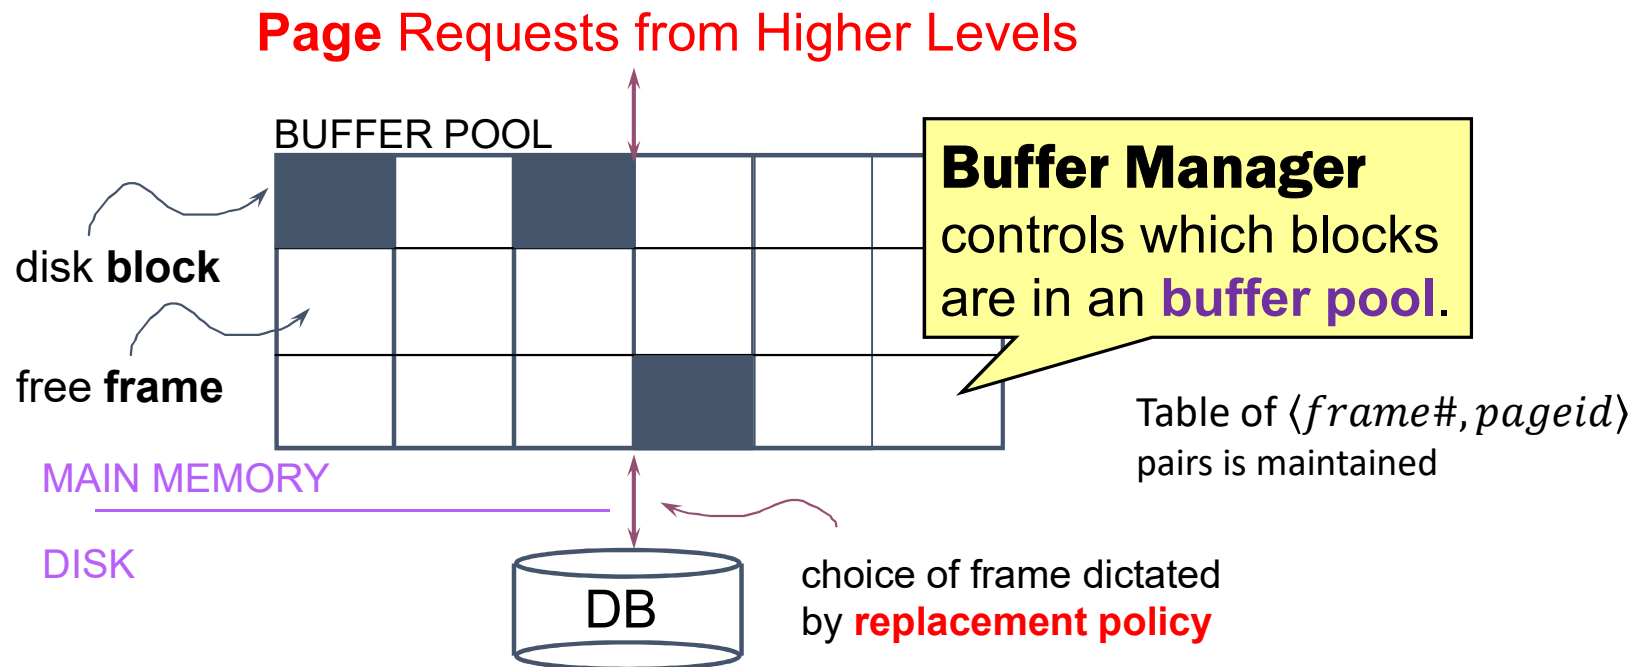
More detail in some other class ...

Store in DISK, Process in (faster) MEMORY

Data must be brought into memory to be processed!

- A requested page is transferred from disk to main memory
(if not fetched earlier & found in memory)
- Evicted pages are transferred from memory to disk
- I/O time dominates the time taken for database operations!
- To minimize I/O time, it is necessary to store and locate data strategically

Disk Buffer Management in a DBMS



- Remember “Store in DISK, Process in (faster) MEMORY”
- Files are moved between disk and main memory in **blocks**; it takes roughly 10 milliseconds
- It is vital that a disk block we are accessing is already in a **buffer pool**!

When a page is requested ...

- A **page** is the unit of memory we request
- If *Page* is in the pool
 - Great! No need to go to disk.
- If not? Choose a frame to replace ...
 - If there is a free frame, use it.
*We **pin** a page = it's in use.*
 - If not? We need to choose a page to remove.

How DBMS makes choice is a **replacement policy**

Goal: minimization of disk access.

Buffer Replacement Policies (cont.)

- **First-In-First-Out (FIFO):**
 - Rule: The oldest arrival page in buffer pool will be replaced first
 - Very simple, intuitive and less maintenance than LRU, but perform poorly in practice ☹️
- **Clock-sweep:**
 - Keep circular list of pages in memory. Each with a 0/1 flag.
A *hand* points to the last examined page frame in the list.
If replacement is needed, the flag at the hand's location is inspected.
 - If 0, the new page is put in place, and advanced the hand one position.
 - Otherwise, the flag is set to 0, then the clocked move until a page is replaced.
 - A more efficient version of FIFO (no need to constantly push page to the back of the queue list.)

Buffer Replacement Policies

- **Least recently used (LRU):**
 - Blocks referenced recently are likely to be used again.
 - Maintain a table to indicate the last time the block in each buffer has been accessed
 - Rule: Throw out the block that has not been read or written for the longest time.
- **Most recently used (MRU): ...**

File Organization

Recap: Data Storage Principles

- Database relations are implemented as **files** of **records**.
- The **storage medium** are **disks**, which consist of **pages**
- Pages are *read* from disk and *written* to disk: *high cost operations!*
- Each record has a **record identity** (rid), which identifies the *page* where it is stored and its *offset* on that page
- The DBMS reads (and writes) entire pages and stores a number of them in a **buffer pool**
- The **buffer manager** decides which pages to load into the buffer

File Organization: *issues*

- **File** format: how are the pages organized in a file
- **Page** format: how are records organized in a page
- **Record** formats: how are attributes organized in a record

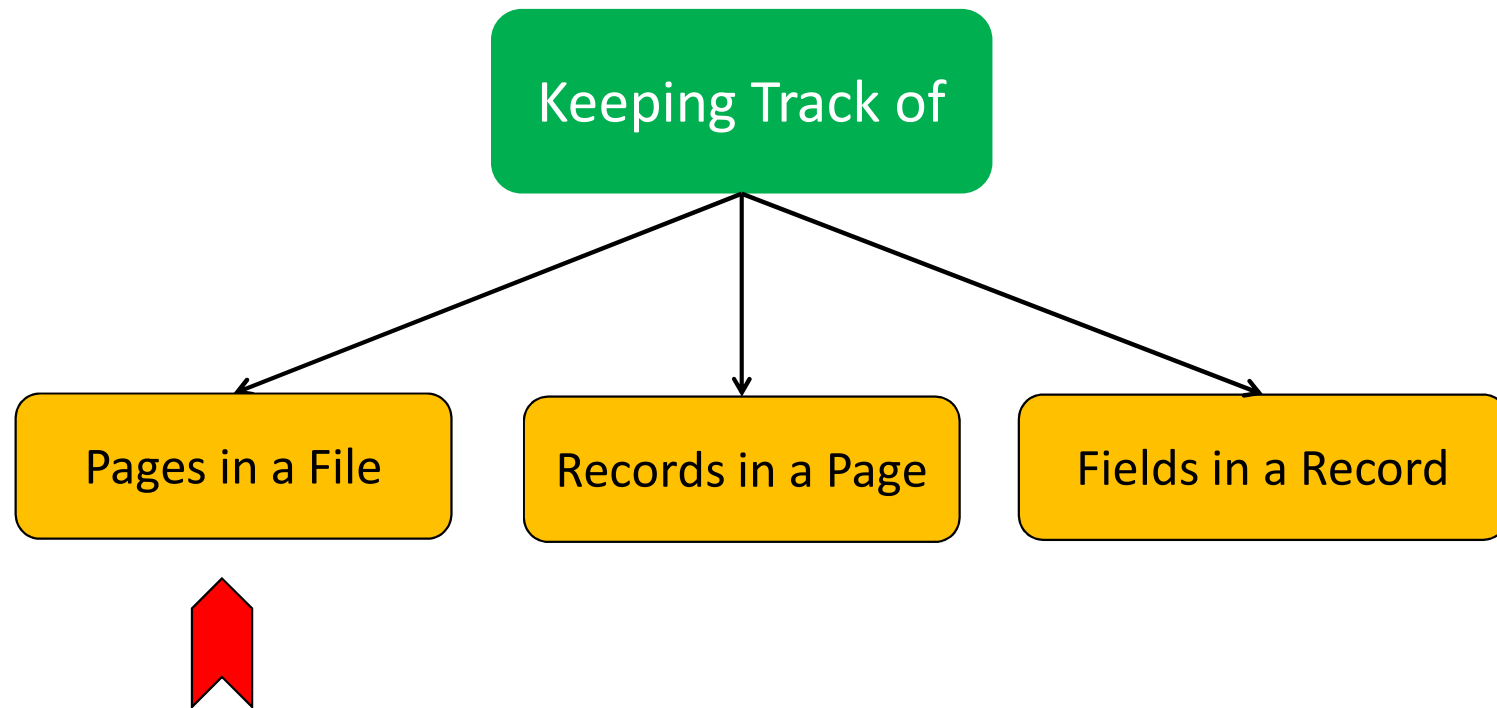
Files must support operations like:

- INSERT/DELETE/MODIFY records
- READ a particular record
- SCAN all records (possibly with some condition)

File Format – Heap Files

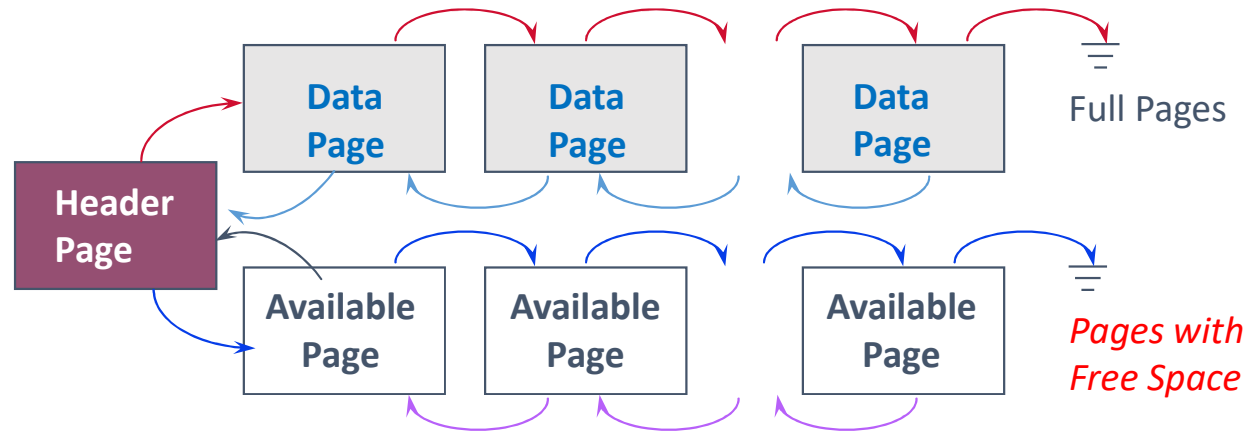
- Data in heap file is **not ordered**.
- Database is stored as a collection of **files**.
- As a heap file grows and shrinks, disk pages are allocated and de-allocated.
- To support record level operations, we must
 - **Keep** track of the **pages** in a file
 - Keep track of the **records** on each page
 - Keep track of the **fields** on each record

Supporting Record-Level Operations



Heap Files using *Linked-Lists* of Pages

A heap file can be organized as a doubly-linked list of pages




The head page, $\langle heap_file_name, page_addr \rangle$, is stored in a known location on disk.
Each page contains 2 pointers plus data

Heap Files using *Linked-Lists* of Pages

- It is likely that page has at least a few free bytes
- Thus, virtually all pages in a file will be on the free list!
- To insert a typical record, we must retrieve and examine several pages on the free list before one with *enough* free space is found.
- This problem can be addressed using an alternative design known as the **directory-based heap file organization**.

Heap Files using *Directory* of Pages

A directory of pages can be maintained whereby each directory entry identifies a page in the heap file



Free space can be managed via maintaining:

- A **bit** per entry (indicating whether the corresponding page has any **free space**)
- A **count** per entry (indicating the amount of **free space** on the page)

Alternative File Formats

Alternatives are ideal for some situation, and not so good in others:

- **Heap Files:**

No order on records. Suitable when typical access is a file scan retrieving all records.

- **Sorted Files:**

Sorted by specific record field (or key).

Best if records must be retrieved in some order, or only a range of records is needed.

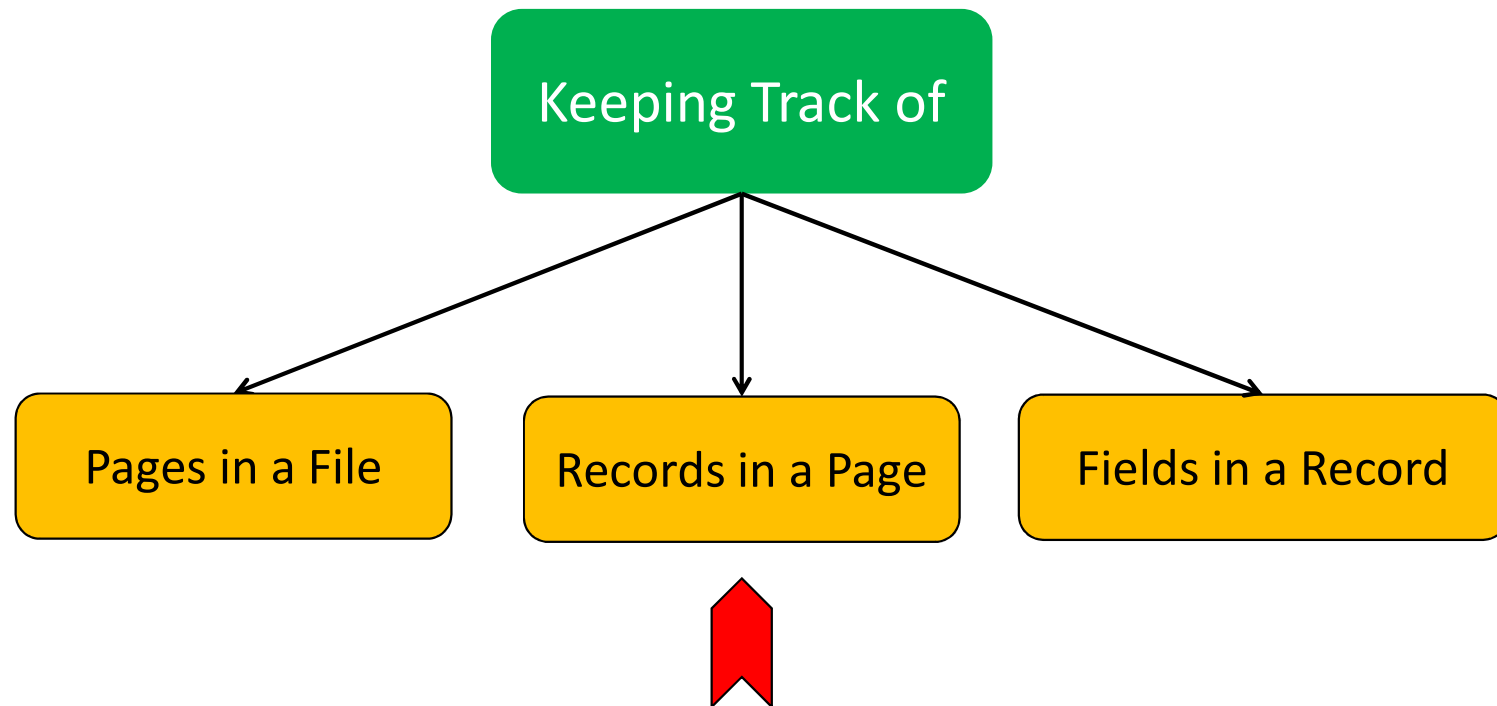
- **Hashed Files:**

File is a collection of **buckets**. (primary page + zero or more overflow pages)

Hash Function h : compute $h(r)$ = bucket into which record r belongs.


Good for equality search (have more trouble with range search).

Supporting Record-Level Operations



Page Formats

A page is a collection of **slots**, which contains records




Each record is identified (commonly) by
$$RID = \langle page_id, slot_number \rangle$$

Issues to consider:

- 1 page = fixed size (e.g., 8KB)
- Records: fixed length, variable length

Fixed-length Records (cont.)


- When records are of fixed-length, slots become *uniform* and can be arranged *consecutively*



- Records can be located by simple offset calculations
- To delete a record, the record is moved into the vacated slot

Fixed-length Records (cont.)

- Alternatively, we can handle deletions by using an array of bits




- When a record is deleted, its bit is turned off.
- Thus, the RIDs of currently stored records remain the same!

Variable-Length Records

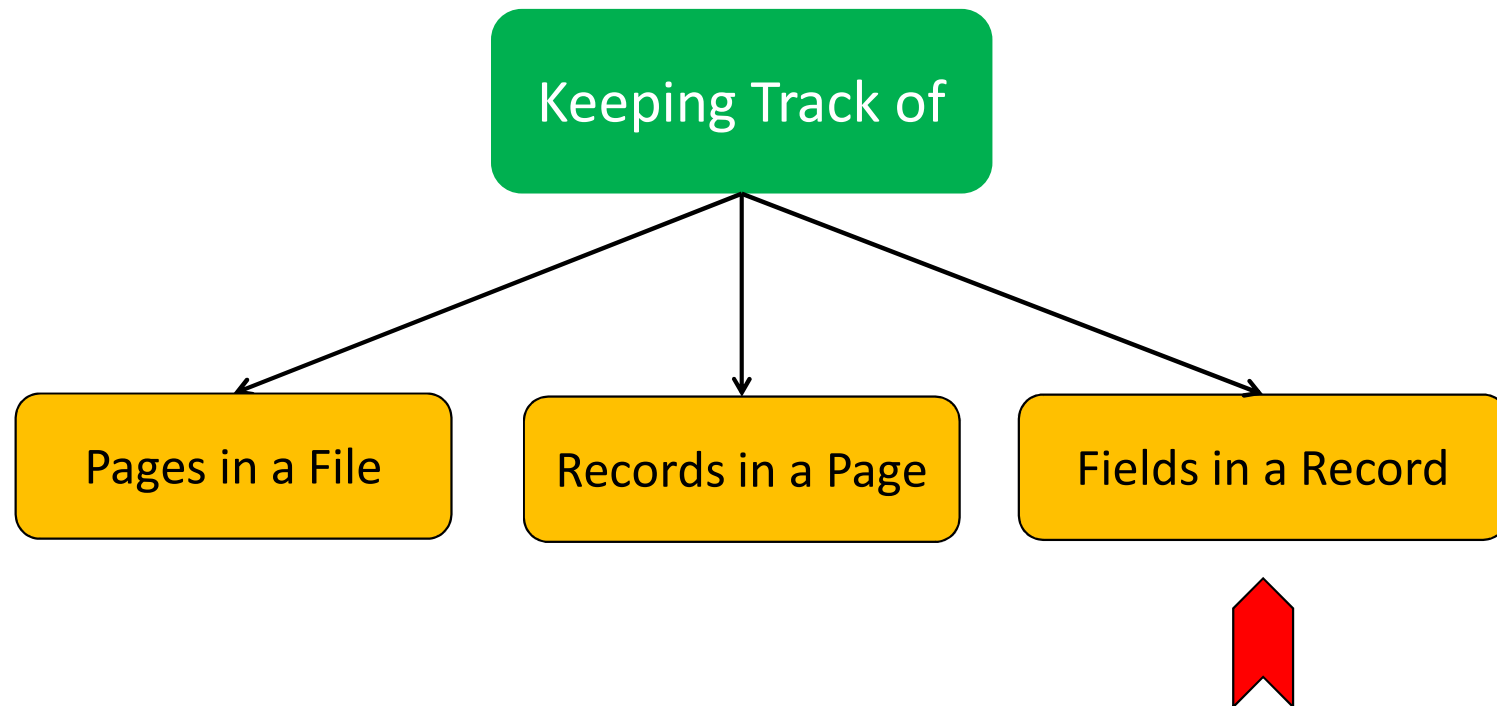
- If the records are of variable length,
we cannot divide the page into a fixed collection of slots
- When a new record is to be inserted,
we have to find an empty slot of “just” the right length
- Thus, when a record is deleted,
we better ensure that all the free space is contiguous
- The ability to move records “without changing RIDs” becomes crucial

Variable-Length Records (cont.)

Maintain a kind of **directory of slots** $\langle record_offset, record_length \rangle$ for each page



Supporting Record-Level Operations




Record Formats

Fields in a record can be either of:

- **Fixed-Length**: each field has a fixed length and the number of fields is also fixed
 - **Variable-Length**: fields are of variable lengths but the number of fields is fixed
-
- Information common to all records (e.g., number of fields and field types) are stored in the system catalog

Fixed-Length Fields

Fixed-length fields can be stored consecutively and their addresses can be calculated using information about the lengths of preceding fields



Variable-Length Fields

There are two possible organizations to store variable-length fields

1. Consecutive storage of fields **separated by delimiters**




This entails a scan of records to locate a desired field!

Variable-Length Fields

There are two possible organizations to store variable-length fields

1. Consecutive storage of fields separated by delimiters
2. Storage of fields with an array of integer offsets



*This offers **direct access** to a field in a record more efficiently!*

Representing Data

How to lay out a tuple (=record)

```
CREATE TABLE Product (  
    pid INT PRIMARY KEY,  
    name CHAR(20),  
    wholesale BIT,  
    description VARCHAR(200));
```


pid 4 B	name 21 B	wholesale 1 bit	description 200 B
------------	--------------	--------------------	----------------------

First guess

How to lay out a tuple (=record)

```
CREATE TABLE Product (  
  pid INT PRIMARY KEY,  
  name CHAR(20),  
  wholesale BIT,  
  description VARCHAR(200));
```

because it is too slow to
parse things that don't
align with boundaries



Second guess

How to lay out a DB page (= block)

DB page/block = multiple of disk block size


In practice, 8 KB or more



How to lay out a fixed-length records

We know neither the length of each record
or the size of each field in it

First attempt



How to lay out a fixed-length records

Second attempt

Block header: schema, length, timestamp
tuple/record
tuple/record
tuple/record
tuple/record
<i>free space</i>

How to lay out a (fixed-length) records

Variable-length records can be handle using offset + record length (as mentioned earlier)

Second attempt (with detail)

Block header: schema, length, timestamp					
pid	name		wholesale	description	
pid	name		wholesale	description	
pid	name		wholesale	description	
pid	name		wholesale	description	
free space					

What about tuples bigger than a page?

spanned tuples

page header	(offset1, length1)	
tuple 1		
page header	(offset1, length1)	
	length2)	
tuple 2		
tuple 1		

You should **seriously** consider changing the DB page size.